

MATH 174 - Numerical Analysis

Dani Nguyen

Fall 2025

Contents

1	Error	2
1.1	Approximation	2
1.2	Round-off Error	2
1.3	Floating Point Arithmetic	3
1.3.1	Degrading Error	3
1.3.2	Catastrophic Cancellation	4
2	Speed	4
2.1	Dot Product	4
2.1.1	MATLAB Dot Product Function	5
2.2	Memory	5
3	Gaussian Elimination	6
3.1	Triangular Matrices	6
3.1.1	Speed of Back Substitution	8
3.2	Augmented Matrices	8
3.2.1	LU Factorization	8
3.3	Matrix Inverses	9
3.3.1	MATLAB Back Substitution Function	9
3.3.2	Gaussian Elimination FLOP count	9
4	Gaussian Elimination With Pivoting	10
4.1	Error	11
4.1.1	Section 7.5	11
4.2	Row Swapping	12

4.3	Speed	12
5	Approximations	13
5.1	Fixed Point Problems	13
5.1.1	Jacobi Iterative Method	14
5.1.2	Gauss-Seidel Iterative Method	14
5.2	Convergence	14

1 Error

1.1 Approximation

Let x be an exact number.

Let y be an approximation of x .

Def: The absolute error is $|x - y|$.

Def: The relative error is $\frac{|x - y|}{|x|}$,

1.2 Round-off Error

Usually a machine uses a finite amount of memory to store a number.

machine $\leftarrow \pi$

Suppose a machine can store some number of digits.

1. Allocate a digit for sign, positive or negative.
2. Allocate a number of digits for the mantissa.
3. Allocate a digit for sign of the exponent.
4. Allocate a number of digits for the exponent.

Under rounding in the case of a 4 digit mantissa, $\pi = 0.3142 \cdot 10^1$, this is the machine # representation of pi.

Def: Round-off error are all errors arising from the error between actual # and machine #.

Notation:

$$x = \text{real \#}$$
$$fl(x) = \text{machine \# representation of } x$$

In an s-digit rounding machine,

$$\text{Absolute error} = |x - fl(x)|$$

$$\text{Relative error} = \frac{|x - fl(x)|}{|x|} \leq \frac{1}{2} \cdot 10^{1-5} = 5 \cdot 10^{-5}$$

1.3 Floating Point Arithmetic

What happens under arithmetic operations? (+, -, ·, /)

Two things to watch out for:

1. Many operations can degrade error.
2. Catastrophic cancellation.

1.3.1 Degrading Error

number	machine
x	$fl(x)$
y	$fl(y)$
$x + y$	$fl(fl(x) + fl(y))$

The more nested $fl()$ there are, the larger the error will be. The same follows for -, ·, /.

When performing arithmetic, it temporarily allows for more digits than the mantissa allows.

Algorithmic "solutions"

$$((x + y) + z) + w$$

Worst case, 4 times the error.

$$(x + y) + (z + w)$$

Worst case, 3 times the error.

1.3.2 Catastrophic Cancellation

Ex: 3-digit rounding machine. Given 2 numbers: 0.312 and 0.311.

$$0.312 - 0.311 = 0.001 = 0.1 \cdot 10^{-2}$$

Loss of significant digits!

Ex: 5-digit rounding machine. Given 2 numbers:

$$x = 0.41626324$$

$$y = 0.41626323$$

$$x - y = 0.00000001$$

In machine:

$$fl(x) = 0.41626$$

$$fl(y) = 0.41626$$

$$fl(x) - fl(y) = 0.00000 = 0$$

$$fl(fl(x) - fl(y)) = 0$$

Relative error:

$$\begin{aligned} \frac{|x - y - fl(fl(x) - fl(y))|}{|x - y|} &= \frac{|x - y - 0|}{|x - y|} \\ &= 1 = 100\% \text{ error!} \end{aligned}$$

2 Speed

2.1 Dot Product

Consider dot product of vectors

$$\vec{x} \in \mathbb{R}^n, \vec{y} \in \mathbb{R}^n$$

$$\begin{aligned} \vec{x} \cdot \vec{y} &= \begin{bmatrix} x_1 & x_2 & \dots & x_n \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \\ &= x_1 y_1 + x_2 y_2 + \dots + x_n y_n \end{aligned}$$

Count time-consuming operations:

1. Add/subtract
2. Multiply/divide
3. Comparisons
4. Memory access/storage

These are called FLOPS (floating point operations).

For the dot product, there are n multiplications and $n-1$ additions.

Adding these together, we get $2n - 1$ FLOPS.

Note that this expression is linear in n .

For large n , $2n$ dominates, $= \mathcal{O}(n) \approx \text{constant times } n$.

Ex: For n vs $2n$, $2n$ will take twice as long.

Ex: For $\mathcal{O}(n^2)$, n vs $2n$, $2n$ will take four times as long.

2.1.1 MATLAB Dot Product Function

A simple MATLAB function to compute the dot product of two vectors and return the number of FLOPS used.

```
function [thedotproduct, theflops] = dotprod(x, y, n)
    thedotproduct = x(1)*y(1);
    theflops = 1;
    for i = 2:n
        thedotproduct = thedotproduct + x(i)*y(i);
        theflops = theflops + 2;
    end
end
```

2.2 Memory

Consider a matrix $A_{n \times n}$

$$\begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} \end{bmatrix} v s \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

The square matrix uses n^2 memory locations, while the vector uses n memory locations. Each number also uses 8 bytes, creating another goal for algorithm development, reduce memory usage.

3 Gaussian Elimination

3.1 Triangular Matrices

Write a MATLAB function that when given $A\vec{x} = \vec{b}$, computes $\vec{x}$

$$A\vec{x} = \vec{b}$$

A is triangular and invertible

Let's say A is lower triangular. $a_{ij} = 0, \forall j > i$

$$\begin{bmatrix} a_{11} & 0 & 0 & \dots & 0 \\ a_{21} & a_{22} & 0 & \dots & 0 \\ a_{31} & a_{32} & a_{33} & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & a_{n3} & \dots & a_{nn} \end{bmatrix}$$

Similarly, if A is upper triangular, $a_{ij} = 0, \forall j < i$.

Ex: $A\vec{x} = \vec{b}$, A is upper triangular.

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ 0 & a_{22} & a_{23} \\ 0 & 0 & a_{33} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

This gives us the following system of equations:

$$\begin{aligned} a_{11}x_1 + a_{12}x_2 + a_{13}x_3 &= b_1 \\ a_{22}x_2 + a_{23}x_3 &= b_2 \\ a_{33}x_3 &= b_3 \end{aligned}$$

So $x_3 = \frac{b_3}{a_{33}}, a_{33} \neq 0$ if A is invertible.

Then, $x_2 = \frac{b_2 - a_{23}x_3}{a_{22}}, a_{22} \neq 0$ if A is invertible.

Finally, $x_1 = \frac{b_1 - a_{12}x_2 - a_{13}x_3}{a_{11}}$, $a_{11} \neq 0$ if A is invertible.

This is called back substitution.

For $i = n, n-1, \dots, 1$

$$\begin{aligned} x_i &= \frac{b_i - (a_{i,i+1}x_{i+1} + \dots + a_{i,n}x_n)}{a_{ii}} \\ &= \frac{b_i - \sum_{j=i+1}^n a_{ij}x_j}{a_{ii}} \end{aligned}$$

Ex: $A\vec{x} = \vec{b}$, A is lower triangular.

$$\begin{bmatrix} a_{11} & 0 & 0 \\ a_{21} & a_{22} & 0 \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

This gives us the following system of equations:

$$\begin{aligned} a_{11}x_1 &= b_1 \\ a_{21}x_1 + a_{22}x_2 &= b_2 \\ a_{31}x_1 + a_{32}x_2 + a_{33}x_3 &= b_3 \end{aligned}$$

This is called forward substitution.

For $i = 1, 2, \dots, n$

$$\begin{aligned} x_i &= \frac{b_i - (a_{i,1}x_1 + \dots + a_{i,i-1}x_{i-1})}{a_{ii}} \\ &= \frac{b_i - \sum_{j=1}^{i-1} a_{ij}x_j}{a_{ii}} \end{aligned}$$

Now given a regular matrix, make it triangular without changing the solutions. Recall Gaussian elimination.

1. E_i (equation i) replaced by cE_i . Row scaling.
2. E_i swapped with E_j . Row swapping.
3. E_i replaced by $E_i + cE_j$. Row addition.

3.1.1 Speed of Back Substitution

For $i = n, n-1, \dots, 1$

$$x_i = \frac{b_i - \sum_{j=i+1}^n u_{ij}x_j}{u_{ii}}$$

$i = n$	1 flop
$n-1$	3 flops
$\vdots$	$\vdots$
1	$2n-1$ flops

$$\sum_{j=1}^n (2j-1) = 2 \sum_{j=1}^n (j) - \sum_{j=1}^n 1 = 2 \left(\frac{n(n+1)}{2} \right) - n = n^2$$

With a large n , the term n^2 dominates, meaning $\mathcal{O}(n^2)$

3.2 Augmented Matrices

Gaussian Elimination on same matrices and different RHS comes up often in applications and algorithms.

3.2.1 LU Factorization

$$A = LU$$

where L is lower triangular with 1's in the diagonal and U is upper triangular

The important entries of L are called multipliers. These entries can be found with the negative of the coefficients of the steps of Gaussian elimination.

$L\vec{y} = \vec{b}$	get y from forward substitution
$U\vec{x} = \vec{y}$	get x from back substitution

3.3 Matrix Inverses

Matrix-vector multiplication:

$$(A\vec{x})_i = \sum_{j=1}^n a_{ij}x_j$$

which gives $2(n-1)n$ flops total

Backwards substitution uses n^2 flops.

Forward substitution uses $n^2 - n$ flops. Because we're working with L, which has 1's on the diagonal, we save n divisions.

In summary, getting A^{-1} to solve $A\vec{x} = \vec{b}$ is slower. Also, A^{-1} has memory storage issues for certain matrices. For example, a tridiagonal matrix, even though it has zeroes on most of its entries, A^{-1} has no zeroes.

3.3.1 MATLAB Back Substitution Function

```
function [x] = backsub(A, b, n)
    for i = n:-1:1
        x(i,1) = b(i);
        for j = i+1:n
            x(i) = x(i)-A(i,j)*x(j)
        end
        x(i,1) = x(i)/A(i,i)
    end
end
```

3.3.2 Gaussian Elimination FLOP count

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & \dots & a_{1n} \\ a_{21} & a_{22} & a_{23} & \dots & a_{2n} \\ a_{31} & a_{32} & a_{33} & \dots & a_{3n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & a_{n3} & \dots & a_{nn} \end{bmatrix}$$

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & \dots & a_{1n} \\ 0 & a_{22} & a_{23} & \dots & a_{2n} \\ a_{31} & a_{32} & a_{33} & \dots & a_{3n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & a_{n3} & \dots & a_{nn} \end{bmatrix}$$

This step has 1 division, 1 multiplication, 1 addition/subtraction.

In total there are $2n - 1$ total flops.

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & \dots & a_{1n} \\ 0 & a_{22} & a_{23} & \dots & a_{2n} \\ 0 & a_{32} & a_{33} & \dots & a_{3n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & a_{n2} & a_{n3} & \dots & a_{nn} \end{bmatrix}$$

Repeat this for $n - 1$ rows (minus 1 because we are excluding the first row).

We can then repeat this for the inner block matrix.

$$(2n - 1)(n - 1) + (2(n - 1) - 1)((n - 1) - 1) + (2(n - 2) - 1)((n - 2) - 1) \dots$$

$$\sum_{j=1}^n j = \frac{n(n+1)}{2}$$

$$\sum_{j=1}^n j^2 = \frac{n(n+1)(2n+1)}{6}$$

4 Gaussian Elimination With Pivoting

But what about $A = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$ Gaussian elimination fails on first step. Can we use another elementary row operation, namely, row swapping?

In general,

$$\begin{bmatrix} 0 & a_{12} & a_{13} & \dots & a_{1n} \\ a_{21} & a_{22} & a_{23} & \dots & a_{2n} \\ a_{31} & a_{32} & a_{33} & \dots & a_{3n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & a_{n3} & \dots & a_{nn} \end{bmatrix}$$

4.1 Error

Find a nonzero element a_{i1} then perform $E_1 \leftrightarrow E_2$. We have to do this before each elimination step. However, there is error involved.

$$\left[\begin{array}{cc|c} \epsilon & 1 & 1 \\ \times & \times & \times \end{array} \right] \rightarrow \tilde{x}_2 = x_2 + \delta$$

$$\rightarrow \tilde{x}_1 = \frac{1 - \tilde{x}_2}{\epsilon} = \frac{1 - x_2 - \delta}{\epsilon} = \frac{1 - x_2}{\epsilon} - \frac{\delta}{\epsilon}$$

Note that in this situation, choosing an x_1 that is small leads to a small ϵ which leads to a large error. We have to be careful when swapping rows.

$$\left[\begin{array}{cc|c} \epsilon & 1 & 1 \\ 0 & \times & \times \end{array} \right]$$

$$\tilde{x}_2 = x_2 + \delta$$

$$|\delta| = |\tilde{x}_2 - x_2|$$

$$\tilde{x}_1 = x_1 - \frac{\delta}{\epsilon}$$

$$\left| \frac{\delta}{\epsilon} \right| = |\tilde{x}_1 - x_1|$$

4.1.1 Section 7.5

$$A\vec{x} = \vec{b}$$

$$A(\vec{x} + \delta\vec{x}) = \vec{b} + \delta\vec{b}$$

$$A\delta\vec{x} = \delta\vec{b}$$

$$\delta\vec{x} = A^{-1}\delta\vec{b}$$

We can use a norm to measure the size of error vectors.

$$\|\delta\vec{x}\| = \|A^{-1}\delta\vec{b}\| \leq \|A^{-1}\| \|\delta\vec{b}\|$$

$$\frac{\|\delta\vec{x}\|}{\|\vec{x}\|} \leq \|A\| \|A^{-1}\| \frac{\|\delta\vec{b}\|}{\|\vec{b}\|}$$

relative error $\leq$ condition $\# \times$ relative size of change of $\vec{b}$

Ex: Suppose you need to solve accurately $A\vec{x} = \vec{b}$ and someone gives you the approx $\tilde{x}$, is it accurate?

We can use the residual vector.

$$\begin{aligned}\vec{r} &= \vec{b} - A\tilde{x} \\ A\tilde{x} &= \vec{b} - \vec{r}\end{aligned}$$

If A is well-conditioned (small condition number), then $\frac{\|\vec{r}\|}{\|\vec{b}\|}$ is small and $\frac{\|\vec{x} - \tilde{x}\|}{\|\vec{x}\|}$ small.

If A is ill-conditioned (large condition number), then we don't know.

4.2 Row Swapping

$$PA = LU$$

where PA is A with rows swapped. P is the identity matrix with its rows swapped to denote the row swapping operations done on A.

Now solving LU factorization:

$$\begin{aligned}A\vec{x} &= \vec{b} \\ PA\vec{x} &= P\vec{b} \\ LU\vec{x} &= P\vec{b} \\ L\vec{y} &= P\vec{b} \quad \text{forward sub} \\ U\vec{x} &= \vec{y} \quad \text{back sub}\end{aligned}$$

4.3 Speed

In an $n \times n$ matrix, it takes $n - 1$ comparisons for the first step, $n - 2$ for the second step, and so on.

$$\begin{aligned}\sum_{j=1}^{n-1} j &= \frac{(n-1)n}{2} \\ &= \mathcal{O}(n^2)\end{aligned}$$

This is negligible additional amount of work for large n compared to $\mathcal{O}(n^3)$ of Gaussian elimination.

Even safer than row partial pivoting in terms of error:

1. Scaled row partial pivoting:
2. Complete pivoting: row swaps and column swaps to move the largest element to pivot element.

5 Approximations

5.1 Fixed Point Problems

Solve an equation $g(x) = h(x)$

Modify to $f(x) = x$

Now we have the fixed point problem, which we can try to solve using fixed point iterations.

Start with guess $x^{(0)}$

Then iterate $x^{(k+1)} = f(x^{(k)})$

The $x^{(k)}$ are approximations of the fixed point.

We can also write $A\vec{x} = \vec{b}$ as fixed point problem using splitting.

$$A = M - N$$

$$(M - N)\vec{x} = \vec{b}$$

$$M\vec{x} = N\vec{x} + \vec{b}$$

$$\vec{x} = M^{-1}(N\vec{x} + \vec{b})$$

where the RHS is $f(\vec{x})$

Choices for M, N:

Let $A = D - L - U$, where D is diagonal, L is lower triangular, U is upper triangular. We have

1. Jacobi iterative method.
2. Gauss-Seidel iterative method.

5.1.1 Jacobi Iterative Method

$$M = D$$

$$N = L + U$$

$$D\vec{x} = (L + U)\vec{x} + \vec{b}$$

$$x_i^{(k+1)} = \frac{b_i - \sum_{j=1}^{i-1} a_{ij}x_j^{(k)} - \sum_{j=i+1}^n a_{ij}x_j^{(k)}}{a_{ii}}$$

5.1.2 Gauss-Seidel Iterative Method

$$M = D - L$$

$$N = U$$

$$(D - L)\vec{x} = U\vec{x} + \vec{b}$$

$$x_i^{(k+1)} = \frac{b_i - \sum_{j>i} a_{ij}x_j^{(k)} - \sum_{j<i} a_{ij}x_j^{(k+1)}}{a_{ii}}$$

5.2 Convergence

Jacobi & Gauss-Seidel are iterative methods. They start with a guess, which improves their approximations at each step. We want $\vec{x}^{(k)}$ sequence of approximations to converge to the exact solution.

Practically, we need a stopping condition. Here are our choices:

1. Stop after certain number of steps
2. Stop when residual is small, where the residual is $\vec{r} = \vec{b} - A\vec{x}$, $||\vec{r}|| \leq \text{tolerance}$.
3. Stop when the algorithm slows down, $||x^{(k+1)} - x^{(k)}|| \leq \text{tolerance}$.

$$A = M - N$$

$$\text{Jacobi} \quad M = D$$

$$\text{Gauss-Seidel} \quad M = D - L$$

$$A\vec{x} = \vec{b} \rightarrow (M - N)\vec{x} = \vec{b}$$

$$\rightarrow M\vec{x} = N\vec{x} + \vec{b}$$

$$\rightarrow \vec{x} = M^{-1}(N\vec{x} + \vec{b})$$

$$\rightarrow \vec{x}^{(k+1)} = M^{-1}(N\vec{x}^{(k)} + \vec{b})$$

$$\rightarrow \vec{x}^{(k+1)} = \mathbf{T}\vec{x}^{(k)} + \vec{c}$$

where $\mathbf{T}$ is the iteration matrix, $M^{-1}N$

and $\vec{c}$ is $M^{-1}\vec{b}$

An iterative method using $\vec{x}^{(k+1)} = \mathbf{T}\vec{x}^{(k)} + \vec{c}$ converges if $\varphi(\mathbf{T}) < 1$, where $\varphi(\mathbf{T})$ is the spectral radius of $\mathbf{T}$.

Def: Spectral radius $\varphi(A) = \max\{|\lambda_1|, |\lambda_2|, \dots, |\lambda_n|\}$

Note: For a general A,